In [3]:
```
import pandas as pd
```

That line of code is to import Pandas for me to be able to use his functions to explore, manipulate, analyze and clean the dataset.

In [4]:
```
df = pd.read_csv('Student_Performance.csv')
```

Pandas function read_csv() is to load the dataset in Jupyter Notebook.

In [5]:
```
duplicates = df.duplicated()
print(duplicates)
```

```
0       False
1       False
2       False
3       False
4       False
        ...
9995    False
9996    False
9997    False
9998    False
9999    False
Length: 10000, dtype: bool
```

As we can see, there are no duplicates in the Dataset.

In [6]:
```
miss_val = df.isnull().any(axis=1)
print(miss_val)
```

```
0       False
1       False
2       False
3       False
4       False
        ...
9995    False
9996    False
9997    False
9998    False
9999    False
Length: 10000, dtype: bool
```

In [11]:
```
df['Extracurricular Activities'] = df['Extracurricular Activities'].replace({'Yes':1, 'No':0})
```

As we can see, there is no row with a field column missing value.

In [12]:
```
import numpy as np
import sklearn
from sklearn.linear_model import LinearRegression
```

#MULTIPLE LINEAR REGRESSION

I'm importing numpy library because, in the implementation of the multiple linear regression, I'll need it to perform mathematical operations such as vector operations and matrix manipulations. Moreover, the dependent variable and independent variables are represented as NumPy arrays in scikit-learn and Pandas, so we need to use NumPy to operate on these arrays within our model.

Now I'll be selecting the independent variables also called features and the dependent variable(target varaiable).

In [13]:
```
X = df[['Hours Studied','Previous Scores','Extracurricular Activities','Sleep Hours']]
Y = df['Performance Index']
```

In [ ]:
```

```

Create and fit the linear regression model

In [14]:
```
model = LinearRegression()
model.fit(X, Y)
```

Out[14]:
```
LinearRegression()
```

Make predictions

In [15]:
```
predictions = model.predict(X)
```

Print model summary (coefficients, R-squared, etc.)

In [17]:
```
print(model.coef_)  # Coefficients
print(model.intercept_)  # Intercept
print(model.score(X, Y))  # R-squared value
```

```
[2.85673907 1.01868958 0.62741988 0.48194429]
-33.24005596528127
0.9879162180086278
```

In [ ]:
```
The coefficients* represent the slope/tilt of the linear relationship between the dependent variable and each one of the
independent variables.
The intercept represent the average expected value for the dependent variable when all independent variables are equal to Zero.
The R-values associated with each coefficient indicate whether the relationship between that predictor variable and the response is statistically significant.
A low p-value (typically < 0.05) suggests a significant relationship, while a high p-value suggests no significant effect.
```

In [ ]:
```
                                        #SIMPLE LINEAR REGRESSION
```

Create the model

In [19]:
```
model = LinearRegression()
```

Fit the model (X: explanatory variable; y: response variable)

In [20]:
```
X = df[['Hours Studied']]
y = df['Sleep Hours']
model.fit(X, y)
```

Out[20]:
```
LinearRegression()
```

Get coefficients (intercept and slope)

In [21]:
```
intercept = model.intercept_
slope = model.coef_[0]

print(f"Intercept (b0): {intercept:.2f}")
print(f"Slope (b1): {slope:.2f}")
```

```
Intercept (b0): 6.53
Slope (b1): 0.00
```

In [ ]:
```

```